

An Empirical Study of Partial Deduction for MINIKANREN

Ekaterina Verbitskaia

JetBrains Research
Saint Petersburg, Russia
kajigor@gmail.com

Daniil Berezun

Saint Petersburg State University
JetBrains Research
Saint Petersburg, Russia
d.berezun@2009.spbu.ru

Dmitry Boulytchev

Saint Petersburg State University
JetBrains Research
Saint Petersburg, Russia
dboulytchev@math.spbu.ru

We study conjunctive partial deduction, an advanced specialization technique aimed at improving the performance of logic programs, in the context of relational programming language MINIKANREN. We identify a number of issues, caused by MINIKANREN peculiarities, and describe a novel approach to specialization based on partial deduction and supercompilation. While the results of the evaluation do not yet demonstrate the advantages of our approach in all cases, we nevertheless consider it as the first step towards an efficient optimization framework for MINIKANREN.

1 Introduction

A family of embedded domain-specific languages MINIKANREN¹ implements relational programming — a paradigm closely related to the pure logic programming. The main primitive in MINIKANREN is *term unification*. Unifications are combined into *conjunctions* and *disjunctions* which form *relations*. Relations can call other relations and be recursive. A program in MINIKANREN consists of a set of relations. To run a MINIKANREN program, one specifies a *goal*: a MINIKANREN expression with some free variables. MINIKANREN then searches for the values of the free variables which satisfy given goal.

The core feature of relational programming is the ability to run a program in various directions by executing goals with free variables. The distribution of free variable occurrences determines the direction of relational search. For example, having specified a relation for adding two numbers, one can also compute the subtraction of two numbers or find all pairs of numbers which can be summed up to get the given one. Program synthesis can be done by running a relational interpreter for some language *backwards*. In general, it is possible to create a solver from a recognizer by translating it into MINIKANREN and running in the appropriate direction [?].

The search employed in MINIKANREN is complete which means that every answer will be found, although it may take a long time. The promise of MINIKANREN falls short when speaking of performance. The execution time of a program in MINIKANREN is highly unpredictable and varies greatly for various directions. What is even worse, it depends on the order of the relation calls within a program. One order can be good for one direction, but slow down the computation dramatically in the other direction.

Specialization or partial evaluation [?] is a technique aimed at improving the performance of a program given some information about it beforehand. It may either be a known value of some argument, its structure (i.e. the length of an input list) or, in case of a relational program, — the direction in which it is intended to be run. An earlier paper [?] has shown that *conjunctive partial deduction* [?] can sometimes improve the performance of MINIKANREN programs. Depending on the particular *control* decisions, it may also not affect the execution time of a program or even make it slower.

Control issues in partial deduction of logic programming language PROLOG have been studied before [?]. Unlike Prolog, where atoms in the right-hand side of a clause cannot be arbitrarily reordered

¹MINIKANREN language web site: <http://minikanren.org>

without changing the semantics of a program, in MINIKANREN the subgoals of conjunction/disjunction can be freely switched. This opens a yet another possibility for optimization, not taken into account by approaches initially developed in the context of conventional logic programming.

In this paper, we study issues which conjunctive partial deduction faces being applied for MINIKANREN. We also describe a novel approach to partial deduction for relational programming, *conservative partial deduction*. We implemented this approach and compared it with the existing specialization system (ECCE) for several programs. We report here the results of the comparison and discuss why some MINIKANREN programs run slower after specialization.

2 Related Work

Specialization is an attractive technique aimed to improve the performance of a program if some of its arguments are known statically. Specialization is studied for functional, imperative, and logic programming and comes in different forms: partial evaluation [?], and partial deduction [?], supercompilation [?], distillation [?], and many others.

The heart of supercompilation-based techniques is *driving* — a symbolic execution of a program through all possible execution paths. The result of driving is so-called *process tree* where nodes correspond to *configurations* which represent computation state. For example, in the case of pure functional programming languages, the computational state might be a term. Each path in the tree corresponds to some concrete program execution. The two main sources for supercompilation optimizations are aggressive information propagation about variables' values, equalities and disequalities, and precomputing of all deterministic semantic evaluation steps. The latter process, also known as *deforestation* [?], means combining of consecutive process tree nodes with no branching. When the tree is constructed, the resulting, or *residual*, program can be extracted from the process tree by the process called *residualization*. Of course, process tree can contain infinite branches. *Whistles* — heuristics to identify possibly infinite branches — are used to ensure supercompilation termination. If a whistle signals during the construction of some branch, then something should be done to ensure termination. The most common approaches are either to stop driving the infinite branch completely (no specialization is done in this case and the source code is blindly copied into the residual program) or to fold the process tree to a *process graph*. The main instrument to perform such a folding is *generalization*. Generalization, abstracting away some computed data about the current term, makes folding possible. One source of infinite branches is consecutive recursive calls to the same function with an accumulating parameter: by unfolding such a call further one can only increase the term size which leads to nontermination. The accumulating parameter can be removed by replacing the call with its generalization. There are several ways to ensure process correctness and termination, most-specific generalization (anti-unification) and *homeomorphic embedding* [?, ?] as a whistle being the most common.

While supercompilation generally improves the behaviour of input programs and distillation can even provide superlinear speedup, there are no ways to predict the effect of specialization on a given program in general. What is worse, the efficiency of residual program from the target language evaluator point of view is rarely considered in the literature. The main optimization source is computing in advance all possible intermediate and statically-known semantics steps at program transformation-time. Other criteria, like the size of the generated program or possible optimizations and execution cost of different language constructions by the target language evaluator, are usually out of consideration [?]. It is known that supercompilation may adversely affect GHC optimizations yielding standalone compilation more powerful [?, ?] and cause code explosion [?]. Moreover, it may be hard to predict the real speedup

of any given program using concrete benchmarks even disregarding the problems above because of the complexity of the transformation algorithm. The worst-case for partial evaluation is when all static variables are used in a dynamic context, and there is some advice on how to implement a partial evaluator as well as a target program so that specialization indeed improves its performance [?, ?]. There is a lack of research in determining the classes of programs which transformers would definitely speed up.

Conjunctive partial deduction [?] makes an effort to provide reasonable control for the left-to-right evaluation strategy of PROLOG. CPD constructs a tree which models goal evaluation and is similar to a SLDNF tree, then a residual program is generated from the tree. Partial deduction itself resembles driving in supercompilation [?]. The specialization is done in two levels of control: the local control determines the shape of the residual programs, while the global control ensures that every relation which can be called in the residual program is defined. The leaves of local control trees become nodes of the global control tree. CPD analyses these nodes at the global level and runs local control for all those which are new.

At the local level, CPD examines a conjunction of atoms by considering each atom one-by-one from left to right. An atom is *unfolded* if it is deemed safe, i.e. a whistle based on homeomorphic embedding does not signal for the atom. When an atom is unfolded, a clause whose head can be unified with the atom is found, and a new node is added into the tree where the atom in the conjunction is replaced with the body of that clause. If there is more than one suitable head, then several branches are added into the tree which corresponds to the disjunction in the residualized program. An adaptation of CPD for the MINIKANREN programming language is described in [?].

ECCE partial deduction system [?] is the most mature implementation of CPD for PROLOG. ECCE provides various implementations of both local and global control as well as several degrees of post-processing. Unfortunately there is neither an automatic procedure to choose what control setting is likely to improve input programs the most nor any informal recommendations on how to choose the best settings. The choice of the proper control is left to the user.

An empirical study has shown that the most well-behaved strategy of local control in CPD for PROLOG is *deterministic unfolding* [?]. An atom is unfolded only if precisely one suitable clause head exists for it with the one exception: it is allowed to unfold an atom non-deterministically once for one local control tree. This means that if a non-deterministic atom is the leftmost one within a conjunction, it is most likely to be unfolded, introducing many new relation calls within the conjunction. We believe this is the core problem of CPD which limits its power when applied to MINIKANREN. The strategy of unfolding atoms from left to right is reasonable in the context of PROLOG because it mimics the way programs in PROLOG execute. But it often leads to larger global control trees and, as a result, bigger, less efficient programs. On the contrary, the results of evaluation of a MINIKANREN program do not depend on the order of atoms (relation calls) within conjunctions, thus we believe a better result can be achieved by selecting a relation call which can restrict the number of branches in the tree. We describe our approach, which implements this idea, in the next section.

3 Conservative Partial Deduction

In this section, we describe a novel approach to relational programs specialization. This approach draws inspiration from both conjunctive partial deduction and supercompilation. The aim was to create a specialization algorithm which would be simpler than conjunctive partial deduction and use properties of MINIKANREN to improve the performance of the input programs.

The algorithm pseudocode is shown in Fig. ???. For the sake of brevity and clarity, we provide

```

1 | conspd goal = residualize ◦ drive ◦ normalize (goal)
2 | drive      = drive_disj ∪ drive_conj
3 |
4 | drive_disj :: Disjunction → Process_Tree
5 | drive_disj D@(c1, ..., cn) =  $\bigvee_{i=1}^n t_i \leftarrow \text{drive\_conj } (c_i)$ 
6 |
7 | drive_conj :: (Conjunction, Substitution) → Process_Tree
8 | drive_conj ((r1, ..., rn), subst) =
9 |   C@(r1, ..., rn) ← propagate_substitution subst on r1, ..., rn
10 |   case whistle (C) of
11 |   | instance (C', subst')      ⇒ create_fold_node (C', subst')
12 |   | embedded_but_not_instance ⇒ create_stop_node (C, subst)
13 |   | otherwise ⇒
14 |   |   | case heuristically_select_a_call (r1, ..., rn) of
15 |   |   | | Just r ⇒
16 |   |   | |   | t ← drive ◦ normalize ◦ unfold (r)
17 |   |   | |   | if trivial ◦ leafs (t)
18 |   |   | |   | then
19 |   |   | |   |   | C' ← propagate_substitution (C \ r, extract_substitution (t))
20 |   |   | |   |   | drive C'[r ↦ extract_calls (t)]
21 |   |   | |   | else
22 |   |   | |   |   | t ∧ drive (C \ r, subst)
23 |   |   | | Nothing ⇒  $\bigwedge_{i=1}^n t_i \leftarrow \text{drive} \circ \text{normalize} \circ \text{unfold } (r_i)$ 

```

Figure 1: Conservative partial deduction pseudo code

functions `drive_disj` and `drive_conj` which describe how to process disjunctions and conjunctions respectively. Driving itself is a trivial combination of the functions provided (line 2).

A driving process creates a process tree, from which a residual program is later created. The process tree is meant to mimic the execution of the input program. The nodes of the process tree include a *configuration* which describes the state of program evaluation at some point. In our case a configuration is a conjunction of relation calls. The substitution computed at each step is also stored in the tree node, although it is not included in the configuration.

Hereafter, we consider all goals and relation bodies to be in *canonical normal form* — a disjunction of conjunctions of either calls or unifications. Moreover, we assume all fresh variables to be introduced into the scope and all unifications to be computed at each step. Those disjuncts in which unifications fail are removed. Each other disjunct takes the form of a possibly empty conjunction of relation calls accompanied with a substitution computed from unifications. Any MINIKANREN term can be trivially transformed into the described form. The function `normalize` in Fig. ?? is assumed to perform term normalization. The code is omitted for brevity.

There are several core ideas behind this algorithm. The first is to select an arbitrary relation to unfold, not necessarily the leftmost which is safe. The second idea is to use a heuristics which decides if unfolding a relation call can lead to discovery of contradictions between conjuncts which in turn leads to restriction of the answer set at specialization-time (line 14; `heuristically_select_a_call` stands for heuristics combination, see section ?? for details). If those contradictions are found, then they are exposed by considering the conjunction as a whole and replacing the selected relation call with the result of its unfolding thus *joining* the conjunction back together instead of using *split* as in CPD (lines 15–22). Joining instead of splitting is why we call our transformer *conservative* partial deduction. Finally, if the

heuristics fails to select a potentially good call, then the conjunction is split into individual calls which are driven in isolation and are never joined (line 23).

When the heuristics selects a call to unfold (line 15), a process tree is constructed for the selected call *in isolation* (line 16). The leaves of the computed tree are examined. If all leaves are either computed substitutions or are instances of some relations accompanied with non-empty substitutions, then the leaves are collected and each of them replaces the considered call in the root conjunction (lines 19–20). If the selected call does not suit the criteria, the results of its unfolding are not propagated onto other relation calls within the conjunction, instead, the next suitable call is selected (line 22). According to the denotational semantics of MINIKANREN it is safe to compute individual conjuncts in any order, thus it is ok to drive any call and then propagate its results onto the other calls.

This process creates branchings whenever a disjunction is examined (lines 4–5). At each step, we make sure to not drive a conjunction which has been already examined. To do this, we check if the current conjunction is a renaming of any other configuration in the tree (line 11). If it is, then we fold the tree by creating a special node which then is residualized into a call to the corresponding relation.

In this approach, we do not generalize in the same fashion as CPD or supercompilation. Our conjunctions are always split into individual calls and are joined back together only if it is meaningful. If the need for generalization arises, i.e. homeomorphic embedding of conjunctions [?] is detected, then we immediately stop driving this conjunction (line 12). When residualizing such a conjunction, we just generate a conjunction of calls to the input program before specialization.

3.1 Unfolding

Unfolding in our case is done by substitution of some relation call by its body with simultaneous normalization and computation of unifications. The unfolding itself is straightforward; however it is not always clear what to unfold and when to *stop* unfolding. Unfolding in the context of specialization of functional programming languages, as well as inlining in specialization of imperative languages, is usually considered to be safe from the residual program efficiency point of view. It may only lead to code explosion or code duplication which is mostly left to a target program compiler optimization or even is out of consideration at all if a specializer is considered as a standalone tool [?].

Unfortunately, this is not the case for the specialization of a relational programming language. Unlike functional and imperative languages, in logic and relational programming languages unfolding may easily affect the target program's efficiency [?]. Unfolding too much may create extra unifications, which is by itself a costly operation, or even introduce duplicated computations by propagating the results of unfolding onto neighbouring conjuncts.

There is a fine edge between too much unfolding and not enough unfolding. The former is maybe even worse than the latter. We believe that the following heuristics provides a reasonable approach to unfolding control.

3.2 Less Branching Heuristics

This heuristics is aimed at selecting a relation call within a conjunction which is both safe to unfold and may lead to discovering contradictions within the conjunction. An unsafe unfolding leads to an uncontrollable increase of the number of relation calls in a conjunction. It is best to first unfold those relation calls which can be fully computed up to substitutions.

We deem every static (non-recursive) conjunct to be safe because they never lead to growth in the number of conjunctions. Those calls which unfold deterministically, meaning there is only one disjunct

```

1 | heuristically_select_a_call :: Conjunction → Maybe Call
2 | heuristically_select_a_call C = find heuristics C
3 |
4 | heuristics :: Call → Bool
5 | heuristics r = isStatic r || isDeterministic r || isLessBranching r

```

Figure 2: Heuristics selection pseudocode

in the unfolded relation, are also considered to be safe.

Those relation calls which are neither static nor deterministic are examined with what we call the *less-branching* heuristics. It identifies the case when the unfolded relation contains fewer disjuncts than it could possibly have. This means that we found some contradiction, some computations were gotten rid of, and thus the answer set was restricted, which is desirable when unfolding. To compute this heuristics we precompute the maximum possible number of disjuncts in each relation and compare this number with the number of disjuncts when unfolding a concrete relation call. The maximum number of disjuncts is computed by unfolding the body of the relation in which all relation calls were replaced by a unification which always succeeds.

The pseudocode describing our heuristics is shown in Fig. ?? . Selecting a good relation call can fail (line 1). The implementation works such that we first select those relation calls which are static, and only if there are none, we proceed to consider deterministic unfoldings and then we search for those which are less branching. We believe this heuristics provides a good balance in unfolding.

4 Evaluation

We implemented the conservative partial deduction for MINIKANREN and compared it with the ECCE partial deduction system. ECCE is designed for PROLOG programming language and cannot be directly applied for programs, written in MINIKANREN. To be able to compare our approach with ECCE, we converted each input program first to the pure subset of PROLOG, then specialized it with ECCE, and then we converted the result back to MINIKANREN. The conversion to PROLOG is a simple syntactic conversion. In the conversion from PROLOG to MINIKANREN, for each Horn clause a conjunction is generated in which unifications are placed before any relation call.

We chose two problems for our study: evaluation of a subset of propositional formulas and type-checking for a simple language. The problems illustrate the approach of using relational interpreters to solve search problems [?]. For both these problems we considered several possible implementations in MINIKANREN which highlight different aspects relevant in specialization.

The `evalo` relation implements an evaluator of a subset of propositional formulas. We consider four different implementations of this relation to explore how the way program is implemented can affect the quality of specialization. Depending on the implementation, ECCE generates programs of varying performance, while the execution times of the programs generated by our approach are similar.

The `typechecko` relation implements a typechecker for a tiny expression language. We consider two different implementations of this relation: one written by hand and the other generated from the functional program. We demonstrate how much these implementations differ in terms of performance before and after specialization.

In this study we measured the execution time for the sample queries, averaging them over multiple runs. We also measured the number of unifications done in search of each individual answer. All ex-

```

let rec evalo subst fm res = conde [fresh (x y z v w) (
  (fm ≡ conj x y ∧ evalo st x v ∧ evalo st y w ∧ ando v w res);
  (fm ≡ disj x y ∧ evalo st x v ∧ evalo st y w ∧ oro v w res);
  (fm ≡ neg x      ∧ evalo st x v ∧ noto v res));
  (fm ≡ var v      ∧ elemo subst v res)]

```

Listing 1: Evaluator of formulas with boolean operation last

```

let rec evalo subst fm res = conde [fresh (x y z v w) (
  (fm ≡ conj x y ∧ ando v w res ∧ evalo st x v ∧ evalo st y w);
  (fm ≡ disj x y ∧ oro v w res ∧ evalo st x v ∧ evalo st y w);
  (fm ≡ neg x      ∧ noto v res  ∧ evalo st x v);
  (fm ≡ var v      ∧ elemo subst v res))]

```

Listing 2: Evaluator of formulas with boolean operation second

amples of MINIKANREN relations in this paper are written in OCANREN². The queries were run on a laptop running Ubuntu 18.04 with quad core Intel Core i5 2.30GHz CPU and 8 GB of RAM.

The tables and graphs use the following denotations. *Original* represents the execution time of a program before any transformations were applied; *ECCE* — of the program specialized by ECCE with default conjunctive control setting; *ConsPD* — of the program specialized by our approach.

4.1 Evaluator of Logic Formulas

The relation eval^o describes an evaluation of a propositional formula under given variable assignments. The relation has three arguments. The first argument is a list of boolean values which plays a role of variable assignments. The i -th value of the substitution is the value of the i -th variable. The second argument is a formula with the following abstract syntax. A formula is either a *variable* represented with a Peano number, a *negation* of a formula, a *conjunction* of two formulas or a *disjunction* of two formulas. The third argument is the value of the formula under the given assignment.

We specialize the eval^o relation to synthesize formulas which evaluate to $\uparrow\text{true}$. To do so, we run the specializer for the goal with the last argument fixed to $\uparrow\text{true}$, while the first two arguments remain free variables. Depending on the way the eval^o is implemented, different specializers generate significantly different residual programs.

4.1.1 The Order of Relation Calls

One possible implementation of the evaluator is presented in Listing ???. Here relation $\text{elem}^o \text{ subst } v \text{ res}$ unifies res with the value of the variable v in the list subst . The relations and^o , or^o , and not^o encode corresponding boolean connectives.

Note, the calls to boolean relations and^o , or^o , and not^o are placed last within each conjunction. This poses a challenge for the CPD-based specializers such as ECCE. Conjunctive partial deduction unfolds relation calls from left to right, so when specializing this relation for running backwards (i.e. considering the goal $\text{eval}^o \text{ subst } fm \uparrow\text{true}$), it fails to propagate the direction data onto recursive calls of eval^o .

²OCANREN: statically typed MINIKANREN embedding in OCAML. The repository of the project: <https://github.com/JetBrains-Research/OCanren>

```

let noto x y = conde [
  (x ≡ ↑true ∧ y ≡ ↑false ;
   x ≡ ↑false ∧ y ≡ ↑true)]

```

Listing 3: Implementation of boolean **not** as a table

```

let noto x y = nando x x y
let oro x y z = nando x x xx ∧ nando y y yy ∧ nando xx yy z
let ando x y z = nando x y xy ∧ nando xy xy z
let nando a b c = conde [
  ( a ≡ ↑false ∧ b ≡ ↑false ∧ c ≡ ↑true );
  ( a ≡ ↑false ∧ b ≡ ↑true ∧ c ≡ ↑true );
  ( a ≡ ↑true ∧ b ≡ ↑false ∧ c ≡ ↑true );
  ( a ≡ ↑true ∧ b ≡ ↑true ∧ c ≡ ↑false )]

```

Listing 4: Implementation of boolean operation via **nand**

Knowing that **res** is ↑**true**, we can conclude that in the call **and**^o v w **res** variables v and w have to be ↑**true** as well. There are three possible options for these variables in the call **or**^o v w **res** and one for the call **not**^o. These variables are used in recursive calls of **eval**^o and thus restrict the result of driving. CPD fails to recognize this, and thus unfolds recursive calls of **eval**^o applied to fresh variables. It leads to over-unfolding, large residual programs and poor performance.


The conservative partial deduction first unfolds those calls which are selected according to the heuristics. Since exploring the implementations of boolean connectives makes more sense, they are unfolded before recursive calls of **eval**^o. The way conservative partial deduction treats this program is the same as it treats the other implementation in which boolean connectives are moved to the left, as shown in Listing ???. This program is easier for ECCE to specialize which demonstrates how unequal the behaviour of CPD for similar programs is.

4.1.2 Unfolding of Complex Relations

Depending on the way a relation is implemented, it may take a different number of driving steps to reach the point when any useful information is derived through its unfolding. Partial deduction tries to unfold every relation call unless it is unsafe, but not all relation calls serve to restrict the search space and thus should be unfolded. In the implementation of **eval**^o boolean connectives can effectively restrict variables within the conjunctions and should be unfolded until they do. But depending on the way they are implemented, the different number of driving steps should be performed for that. The simplest way to implement these relations is by mimicking a truth tables as demonstrated by the implementation of **not**^o in Listing ???. It is enough to unfold such relation calls once to derive useful information about variables.

The other way to implement boolean connectives is to express them using a single basic boolean relation such as **nand**^o which is, in turn, has a table-based implementation (see Listing ??). It will take several sequential unfoldings to derive that variables v and w should be ↑**true** when considering a call **and**^o v w ↑**true** implemented via a basic relation. Conservative partial deduction drives the selected call until it derives useful substitutions for the variables involved while CPD with deterministic unfolding may fail to do so.

	Original	ECCE	ConsPD
<i>First-Plain</i>	1.59s	1.61s	0.92s
<i>First-Nando</i>	1.43s	2.24s	0.96s
<i>Last-Plain</i>	0.98s	1.43s	0.97s
<i>Last-Nando</i>	1.09s	1.54s	0.91s

Figure 3: Execution time of eval^o

4.1.3 Evaluation Results

In our study we considered 4 implementations of eval^o summed up in the table ???. They differ in the way the boolean connectives are implemented (see column *Implementation*) and whether they are placed before or after the recursive calls to eval^o (see column *Placement*). These four implementations are very different from the standpoint of ECCE.

	Implementation	Placement
<i>FirstPlain</i>	table-based	before
<i>LastPlain</i>	table-based	after
<i>FirstNando</i>	via nand ^o	before
<i>LastNando</i>	via nand ^o	after

Table 1: Different implementations of eval^o

We measured the time necessary to generate 1000 formulas over two variables which evaluate to \uparrow true (averaged over 10 runs). The results are presented in Fig. ??.

Conservative partial deduction generates programs with comparable performance for all four implementations, while the quality of ECCE specialization differs significantly. ECCE worsens performance for every implementation as compared to the original program. ConsPD do not worsen performance for any implementation. Its effect is most significant for the implementations in which the boolean connectives are placed first within conjunctions.

4.1.4 The Order of Answers

It is important to note that different implementations of the same MINIKANREN relation produce answers in different orders. Nevertheless, since MINIKANREN search is complete, all answers will be found eventually. Unfortunately, it is not guaranteed that the first 1000 formulas generated with different implementations of eval^o will be the same. For example, 983 formulas are the same among the first 1000 formulas generated by the Original *FirstPlain* relation and the same relation after the ConsPD transformation. At the same time, only 405 formulas are the same between the Original and ECCE *LastNando* relations.

The reason why implementations differ so much in the order of the answers stems from the canonical search strategy employed in MINIKANREN. Most MINIKANREN implementations employ *interleaving* search [?] which is left-biased. It means that the leftmost disjunct in a relation is being executed longer than the disjunct on the right. This property is not local which makes it very hard to estimate the performance of a given relation.

In practice it means that if a specializer reorders disjuncts, then the performance of relations after specialization may be unpredictable. For example, by putting the disjuncts of the `evalo` relation in the opposite order, one produces a relation which runs much faster than the original, but it generates completely different formulas at the same time. Most of the first 1000 formulas in this case are multiple negations of a variable, while the original relation produces more diverse set of answers. Computing a negation of a formula only takes one recursive `evalo` call thus finding such answers is faster than conjunctions and disjunctions. Meanwhile, the formulas generated by the reordered relation are less diverse and may be of less interest.

Although neither ECCE nor ConsPD reorder disjuncts, they remove disjuncts which cannot succeed. Thus they influence the order of answers and performance of relations. We believe that, in general, it is not possible to guarantee the same order of answers after specialization. Exploring how different specializations influence the execution order is a fascinating direction for future research.

4.2 Typechecker-Term Generator

This relation implements a typechecker for a tiny expression language. Being executed in the backward direction it serves as a generator of terms of the given type. The abstract syntax of the language is presented below. The variables are represented with de Bruijn indices, thus let-binding does not specify which variable is being bound.

$$\begin{array}{l}
 \text{type term} = \quad B\text{Const of Bool} \quad | \quad I\text{Const of Int} \quad | \quad \text{Var of Int} \\
 \quad | \quad \text{term} + \text{term} \quad | \quad \text{term} * \text{term} \quad | \quad \text{term} = \text{term} \quad | \quad \text{term} < \text{term} \\
 \quad | \quad \text{let } \underline{\text{term}} \text{ in } \underline{\text{term}} \quad | \quad \text{if } \underline{\text{term}} \text{ then } \underline{\text{term}} \text{ else } \underline{\text{term}}
 \end{array}$$

The typing rules are straightforward and are presented below. Boolean and integer constants have the corresponding types regardless of the environment. Only terms of type integer can be summed up, multiplied or compared by less-than operator. Any terms of the same type can be checked for equality. Addition and multiplication of two terms of suitable types have integer type, while comparisons have boolean type. If-then-else expression typechecks only if its condition is of type boolean, while both then- and else-branches have the same type. An environment Γ is an ordered list, in which the i -th element is the type of the variable with the i -th de Bruijn index. To typecheck a let-binding, first, the term being bound is typechecked and is added in the beginning of the environment Γ , and then the body is typechecked in the context of the new environment. Typechecking a variable with the index i boils down to getting an i -th element of the list.

We compared two implementations of these typing rules. The first one is obtained by unnesting of a functional program as described in [?] (*Generated*). It is worth noting that the unnesting introduces a lot of redundancy in form of extra unifications and thus creates programs which are very inefficient. Thus we contrast this implementation with the program hand-written in OCanren (*Hand-written*). Each implementation has been specialized with ConsPD and ECCE. We measured the time needed to generate 1000 closed terms of type integer (see Fig. ??).

As expected, the generated program is far slower than the hand-written. The principal difference between these two implementations is that the generated program contains a certain redundancy introduced

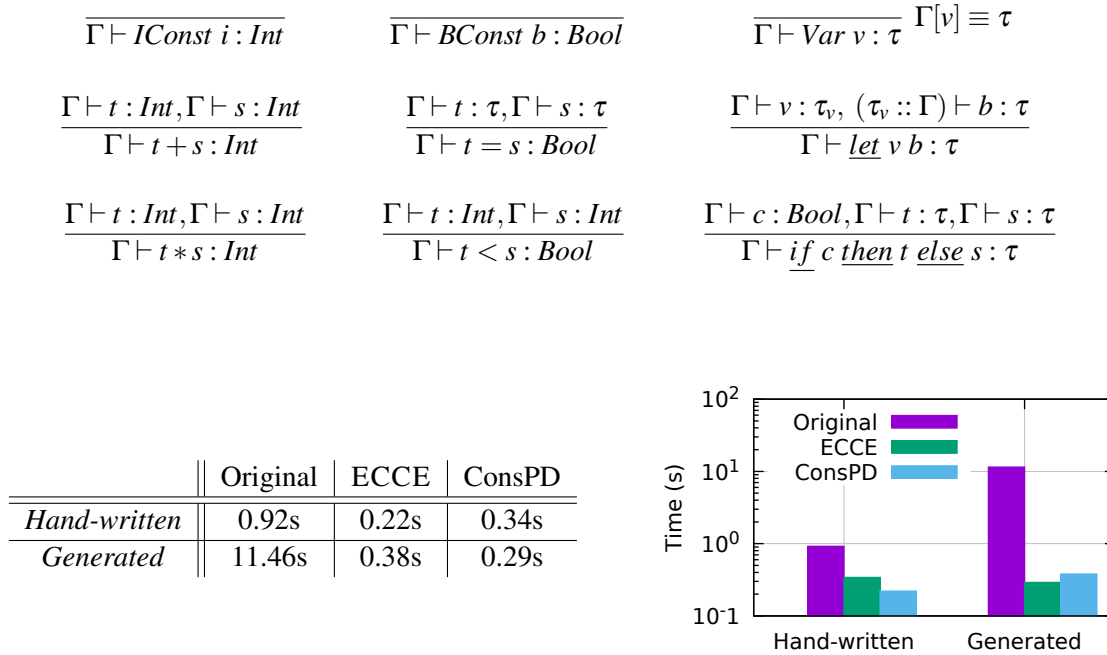


Figure 4: Execution time of generating 1000 closed terms of type Integer

by unnesting. For example, typechecking of the sum of two terms in the hand-written implementation consists of a single conjunction (see Listing ??) while the generated program is far more complicated and also uses a special relation typeEq^o to compare types (see Listing ??).

```

let rec typechecko gamma term res = conde [
  ...
  fresh (x y) ((term ≡ x + y ∧
    typechecko gamma x ↑(Some Integer) ∧
    typechecko gamma y ↑(Some Integer) ∧
    res ≡ ↑(Some Integer)));
  ...]

```

Listing 5: A fragment of hand-written typechecker

Most of the redundancy of the generated program is removed by specialization with respect to the known type of the term. This is why both implementations have comparable speed after specialization. ECCE shows bigger speedup for the hand-written program than ConsPD and vice versa for the generated implementation. We believe that this difference can be explained by too much unfolding. ECCE performs a lot of excessive unfolding for the generated program and only barely changes the hand-written program. At the same time ConsPD specializes both implementations to comparable programs performing average amount of unfolding. This shows that the heuristics we presented gives more stable, although not the best, results.

```

let rec typechecko gamma term res = conde [
  ...
  fresh (x y t1 t2) ((term ≡ x + y ∧
    conde [
      typechecko gamma x ↑None      ∧ res ≡ ↑None;
      typechecko gamma x ↑(Some t1) ∧
        typechecko gamma y ↑None      ∧ res ≡ ↑None;
      typechecko gamma x ↑(Some t1) ∧ typechecko gamma y ↑(Some t2) ∧
        typeEqo t1 Integer ↑true      ∧ typeEqo t2 Integer ↑true ∧
        res ≡ ↑(Some Integer);
    ])
  ...]

```

Listing 6: A fragment of generated typechecker

5 Conclusion

In this paper we discussed some issues which arise in the area of partial deduction techniques for relational programming language MINIKANREN. We presented a novel approach to partial deduction — conservative partial deduction — which uses a heuristics to select a suitable relation call to unfold at each step of driving. We compared this approach with the most sophisticated implementation of conjunctive partial deduction — ECCE partial deduction system — on 6 relations which solve 2 different problems.

Our specializer improved the execution time of all queries. ECCE worsened the performance of all 4 implementations of the propositional evaluator relation, while improving the other queries. Conservative PD is more stable with regards to the order of relation calls than ECCE which is demonstrated by the similar performance of all 4 implementations of the evaluator of logic formulas.

Some queries to the same relation were improved better by ConsPD, while others — by ECCE. We conclude that there is still no one good technique which definitively speeds up every relational program. More research is needed to develop models capable of predicting the performance of a relation which can be used in specialization. Another direction for future research is exploring how specialization influences the execution order of a MINIKANREN program.

References

- [1] Maximilian C. Bolingbroke & Simon L. Peyton Jones (2010): *Supercompilation by evaluation*. In Jeremy Gibbons, editor: *Proceedings of the 3rd ACM SIGPLAN Symposium on Haskell, Haskell 2010, Baltimore, MD, USA, 30 September 2010*, ACM, pp. 135–146, doi:10.1145/1863523.1863540. Available at <https://doi.org/10.1145/1863523.1863540>.
- [2] Mikhail A. Bulyonkov (1984): *Polyvariant Mixed Computation for Analyzer Programs*. *Acta Inf.* 21, pp. 473–484, doi:10.1007/BF00271642. Available at <https://doi.org/10.1007/BF00271642>.
- [3] Danny De Schreye, Robert Glück, Jesper Jørgensen, Michael Leuschel, Bern Martens & Morten Heine Sørensen (1999): *Conjunctive partial deduction: Foundations, control, algorithms, and experiments*. *The Journal of Logic Programming* 41(2-3), pp. 231–277.

- [4] Robert Glück & Morten Heine Sørensen (1994): *Partial deduction and driving are equivalent*. In: *International Symposium on Programming Language Implementation and Logic Programming*, Springer, pp. 165–181.
- [5] Geoff W Hamilton (2007): *Distillation: extracting the essence of programs*. In: *Proceedings of the 2007 ACM SIGPLAN symposium on Partial evaluation and semantics-based program manipulation*, pp. 61–70.
- [6] G. Higman (1952): *Ordering by divisibility in abstract algebras*. In: *Proceedings of the London Mathematical Society*, 2, pp. 326–336.
- [7] Neil D. Jones, Carsten K. Gomard & Peter Sestoft (1993): *Partial evaluation and automatic program generation*. Prentice Hall international series in computer science, Prentice Hall.
- [8] Peter A. Jonsson & Johan Nordlander (2011): *Taming code explosion in supercompilation*. In Siau-Cheng Khoo & Jeremy G. Siek, editors: *Proceedings of the 2011 ACM SIGPLAN Workshop on Partial Evaluation and Program Manipulation, PEPM 2011, Austin, TX, USA, January 24-25, 2011*, ACM, pp. 33–42, doi:10.1145/1929501.1929507. Available at <https://doi.org/10.1145/1929501.1929507>.
- [9] Oleg Kiselyov, Chung-chieh Shan, Daniel P. Friedman & Amr Sabry (2005): *Backtracking, Interleaving, and Terminating Monad Transformers: (Functional Pearl)*. *SIGPLAN Not.* 40(9), p. 192203, doi:10.1145/1090189.1086390. Available at <https://doi.org/10.1145/1090189.1086390>.
- [10] J. B. Kruskal (1960): *Well-quasi ordering, the tree theorem, and Vazsonyi's conjecture*. 95, pp. 210–225.
- [11] Michael Leuschel (1997): *Advanced techniques for logic program specialisation*.
- [12] Michael Leuschel (1997): *The ecce partial deduction system*. In: *Proceedings of the ILPS*, 97, Citeseer.
- [13] Michael Leuschel & Maurice Bruynooghe (2002): *Logic program specialisation through partial deduction: Control issues*. *Theory and Practice of Logic Programming* 2(4-5), pp. 461–515.
- [14] John W. Lloyd & John C Shepherdson (1991): *Partial evaluation in logic programming*. *The Journal of Logic Programming* 11(3-4), pp. 217–242.
- [15] Petr Lozov, Ekaterina Verbitskaia & Dmitry Boulytchev (2019): *Relational Interpreters for Search Problems*. In: *miniKanren and Relational Programming Workshop*, p. 43.
- [16] Neil Mitchell & Colin Runciman (2007): *A Supercompiler for Core Haskell*. In Olaf Chitil, Zoltán Horváth & Viktória Zsók, editors: *Implementation and Application of Functional Languages, 19th International Workshop, IFL 2007, Freiburg, Germany, September 27-29, 2007. Revised Selected Papers, Lecture Notes in Computer Science 5083*, Springer, pp. 147–164, doi:10.1007/978-3-540-85373-2_9. Available at https://doi.org/10.1007/978-3-540-85373-2_9.
- [17] Morten Heine Soerensen, Robert Glück & Neil D. Jones (1996): *A positive supercompiler*. *Journal of functional programming* 6(6), pp. 811–838.
- [18] Philip Wadler (1990): *Deforestation: Transforming Programs to Eliminate Trees*. *Theor. Comput. Sci.* 73(2), pp. 231–248, doi:10.1016/0304-3975(90)90147-A. Available at [https://doi.org/10.1016/0304-3975\(90\)90147-A](https://doi.org/10.1016/0304-3975(90)90147-A).